

Historical Timeline: How Conversational Memory Became Necessary

Deliverable 1 - Problem Reconstruction

Domain: Coding / debugging assistant

How to read this timeline

Each entry is a link in a causal chain, not a standalone summary. An approach is adopted, it hits a specific bottleneck, and that bottleneck forces the next approach. Read top to bottom: each stage exists because the previous one failed at something concrete.

The chain does not name the target solution. It ends at the point where the remaining bottlenecks define what any conversational memory system must address. Those requirements are derived in `first_principles.md`.

The chain

Link 0 - Stateless language models (baseline)

Approach. Send the current input; the model responds from its trained weights, with no memory of prior exchanges.

Bottleneck. Nothing persists between requests. The model cannot use anything the user established earlier.

What it forced. A way to give the model the prior conversation, by placing it in the input.

Link 1 - The Transformer and the fixed context window (2017)

Reference: Vaswani et al., *Attention Is All You Need* [1].

Approach. Process a window of tokens using self-attention, where every token attends to every other token. This made models parallelisable and highly capable, and became the basis of modern LLMs.

Bottleneck. Attention compares every token against every other token [1, §3.2.1], so cost grows as $O(n^2 \cdot d)$ with sequence length [1, Table 1]. The context window is therefore fixed and expensive to extend, and the architecture provides no persistence beyond it.

What it forced. If history cannot be held in the window, and the window cannot be cheaply enlarged, information must be stored outside the model and brought in when needed.

Link 2 - Retrieval-augmented generation (2020)

Reference: Lewis et al., *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks* [2].

Approach. Pair the model (parametric memory) with an external, searchable store (non-parametric memory). Retrieve the top-K passages most similar to the query and generate the answer conditioned on them. Knowledge can be revised, inspected, and attributed, and updated by replacing the store [2, Introduction; §4.5].

Bottleneck. Retrieval ranks only by similarity to the current query, over a shared, general corpus. It does not decide what to store about an individual user, resolve conflicting facts, forget, or separate users. Retrieving more is not reliably better [2, §4.5], and the store is treated as ground truth with no signal when it is stale [2, §4.5]. MemGPT later frames its own focus, long-term memory of user inputs, as distinct from this retrieval-augmented lineage [4, §4].

What it forced. Mechanisms that rank memories on more than similarity, and that decide what is worth keeping and synthesising, not merely what is similar.

Link 3 - Memory ranking and synthesis (April 2023)

Reference: Park et al., *Generative Agents: Interactive Simulacra of Human Behavior* [3].

Approach. Maintain a memory stream and rank memories on three signals, recency, importance, and relevance, rather than similarity alone [3, §4.1]. Periodically synthesise raw observations into higher-level reflections, stored as memories, so the system can generalise [3, §4.2].

Bottleneck. Even with multi-signal ranking, the system suffers retrieval misses, partial-fragment retrieval, and embellishment [3, §6.5.2]; the signal weighting is unsettled [3, §8.2]; and memory can be manipulated by crafted input [3, §8.2]. It does not address multi-user isolation or deletion. Separately, ranking and synthesis do not solve the problem that the underlying context window is still a hard, finite resource.

What it forced. A way to treat the limited context itself as a managed resource, deciding what occupies it at any moment, alongside better ranking.

Link 4 - Context as a managed resource (October 2023)

Reference: Packer et al., *MemGPT: Towards LLMs as Operating Systems* [4].

Approach. Treat the fixed context window as a constrained resource and manage what occupies it, moving information between the window and external storage as needed [4, Introduction]. This reframes the problem from “enlarge the window” to “allocate the window.”

Bottleneck. Context budgets remain finite and vary widely by model [4, Table 1]; summary-based memory answers only about a third of questions about its own prior conversations [4, Table 2]; single-pass retrieval is capped by the retriever and fails on chained questions [4, §3.2.1, §3.2.2]; and the system’s own agent often stops searching before exhausting the store [4, §3.2.1]. Privacy, deletion, and multi-user isolation are not addressed.

What it forced. A dedicated conversational memory layer that answers what to remember, what to forget, how to retrieve, how memory influences the conversation, and how it evolves, while also handling per-user isolation and sensitive data, none of which the prior links resolve.

Where the chain ends

The links above establish, with evidence, that:

- information must live outside the fixed window (Link 1);
- an external store solves knowledge access but not memory management (Link 2);
- memories must be ranked on more than similarity and synthesised, but ranking alone leaves misses, integrity, and isolation unsolved (Link 3);

- the window must be actively managed as a resource, but budgeting alone does not decide what is worth remembering about a user, nor handle forgetting, isolation, or sensitive data (Link 4).

What remains unresolved across all four links defines the conversational memory problem. Those unresolved needs are stated as requirements in `first_principles.md` and framed as open questions in `problem_reconstruction.pdf`.

This progression is corroborated by MemGPT's own account of the field, which identifies three lineages, long-context models, retrieval-augmented models, and LLMs as agents with long-term memory, matching Links 1, 2, and 3–4 respectively [4, §4].

References

- [1] A. Vaswani et al. *Attention Is All You Need*. 2017. <https://arxiv.org/abs/1706.03762>
- [2] P. Lewis et al. *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*. 2020. <https://arxiv.org/abs/2005.11401>
- [3] J. S. Park et al. *Generative Agents: Interactive Simulacra of Human Behavior*. 2023. <https://arxiv.org/abs/2304.03442>
- [4] C. Packer et al. *MemGPT: Towards LLMs as Operating Systems*. 2023. <https://arxiv.org/abs/2310.08560>